

Ch5-ARMA


Fitting ARMA model to lake Huron data

This example shows how to estimate ARMA parameters using R. Recall lake Huron data given by

```
source("TS-library.R")
lake = c(10.38,11.86,10.97,10.8,9.79,
10.39,10.42,10.82,11.4,11.32,11.44,11.68,11.17,
10.53,10.01,9.91,9.14,9.16,9.55,9.67,8.44,
8.24,9.1,9.09,9.35,8.82,9.32,9.01,9,9.8,9.83,9.72,9.89,
10.01,9.37,8.69,8.19,8.67,9.55,8.92,8.09,
9.37,10.13,10.14,9.51,9.24,8.66,8.86,8.05,
7.79,6.75,6.75,7.82,8.64,10.58,9.48,7.38,
6.9,6.94,6.24,6.84,6.85,6.9,7.79,8.18,
7.51,7.23,8.42,9.61,9.05,9.26,9.22,9.38,
9.1,7.95,8.12,9.75,10.85,10.41,9.96,9.61,
8.76,8.18,7.21,7.13,9.1,8.25,7.91,6.89,
5.96,6.8,7.68,8.38,8.52,9.74,9.31,9.89,9.96)
```

From the ACF/PACF plots,

```
par(mfrow=c(1,2))
acf2(lake, lag=30);
pacf(lake)
```



AR(2) seems plausible.

We first consider Yule-Walker estimation of AR(2) model. Note by default, `ar.yw()` includes a constant (non-zero mean model).

```
ar.yw(lake, aic=FALSE, order.max=2)

##
## Call:
## ar.yw.default(x = lake, aic = FALSE, order.max = 2)
##
## Coefficients:
##      1      2
## 1.0538 -0.2668
##
## Order selected 2  sigma^2 estimated as  0.5075
```

```
ar.yw(lake, aic=FALSE, order.max=2, demean=FALSE)

##
## Call:
## ar.yw.default(x = lake, aic = FALSE, order.max = 2, demean = FALSE)
##
## Coefficients:
##      1      2
## 1.0747 -0.0923
##
## Order selected 2  sigma^2 estimated as  2.7
```

If you want to apply MLE to estimate parameters

$$\Phi(B)(X_t - \mu) = \Theta(B)Z_t,$$


do the following

```
ar11.out = arima(lake, order=c(1,0,1), method = c("CSS-ML"), include.mean = TRUE)
ar1.out = arima(lake, order=c(1,0,0))
ar2.out = arima(lake, order=c(2,0,0))
ma1.out = arima(lake, order=c(0,0,1))
```


After fit ARMA, we need to check the goodness of fit using residuals.

```
tsdiag(ar11.out)
```


Standardized Residuals



ACF of Residuals




p values for Ljung–Box statistic




```
tsdiag(ar2.out)
```


Standardized Residuals



ACF of Residuals



p values for Ljung–Box statistic




Test of randomness gives


```
library(itsmr)
test(resid(ar11.out))
```

```
## Null hypothesis: Residuals are iid noise.
## Test Distribution Statistic p-value
## Ljung-Box Q Q ~ chisq(20) 10.14 0.9656
## McLeod-Li Q Q ~ chisq(20) 16.43 0.6899
## Turning points T (T-64)/4.1 ~ N(0,1) 69 0.2266
## Diff signs S (S-48.5)/2.9 ~ N(0,1) 50 0.6015
## Rank P (P-2376.5)/162.9 ~ N(0,1) 2083 0.0716
```


ACF




PACF



Residuals

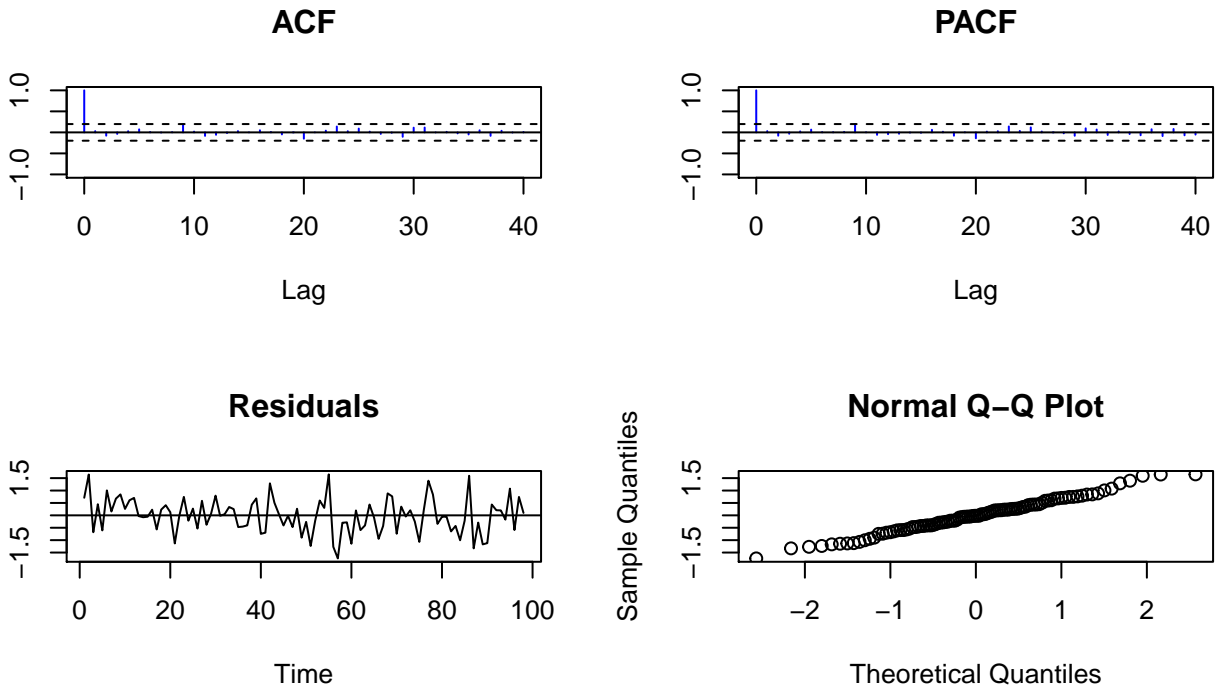


Normal Q-Q Plot



```
test(resid(ar2.out))
```

```
## Null hypothesis: Residuals are iid noise.
## Test Distribution Statistic p-value
## Ljung-Box Q Q ~ chisq(20) 10.67 0.9544
## McLeod-Li Q Q ~ chisq(20) 17.49 0.6211
## Turning points T (T-64)/4.1 ~ N(0,1) 65 0.8089
## Diff signs S (S-48.5)/2.9 ~ N(0,1) 52 0.223
## Rank P (P-2376.5)/162.9 ~ N(0,1) 2051 0.0457 *
```



Both model seems plausible. To find the best model in terms of information criteria,

```
AICC = BIC = AIC = P = Q = NULL;
pmax=3; qmax=3;
n = length(lake);
for(p in 0:qmax){
  for(q in 0:qmax){
    fit = arima(lake, order=c(p, 0, q), include.mean=TRUE);
    m = p+q+2;
    AIC = c(AIC, -2*fit$loglik + 2*m);
    AICC = c(AICC, -2*fit$loglik + 2*m*n/(n-m-1));
    BIC = c(BIC, -2*fit$loglik + m*log(n));
    P = c(P, p);
    Q = c(Q, q);
  }
}
```

```
## Warning in arima(lake, order = c(p, 0, q), include.mean = TRUE): possible
## convergence problem: optim gave code = 1
```

```
id1 = which.min(AICC);
id2 = which.min(BIC);
c(P[id1], Q[id1])
```


```
## [1] 1 1
```

```
c(P[id2], Q[id2])
```


```
## [1] 1 1
```

```
# Both gives ARMA(1,1) as the best model
```

```
plot(AICC);
```



```
plot(BIC);
```



Or simply, you can apply

```
library(forecast)

##
## Attaching package: 'forecast'
## The following object is masked from 'package:itsmr':
##
##      forecast
fit=auto.arima(lake, d=0);
summary(fit)

## Series: lake
## ARIMA(1,0,1) with non-zero mean
##
## Coefficients:
##          ar1      ma1      mean
##          0.7449  0.3206  9.0555
## s.e.    0.0777  0.1135  0.3501
##
## sigma^2 estimated as 0.4899: log likelihood=-103.25
## AIC=214.49  AICc=214.92  BIC=224.83
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.008976993 0.6891588 0.5495562 -0.741374 6.307826 0.9385026
##              ACF1
## Training set 0.004670692
```

We can also perform the test whether all coefficients are away from zero

$$H_0 : \phi_i = 0 \quad vs \quad H_1 : \phi_i \neq 0$$

```
2*(1-pnorm(fit$coef/(sqrt(diag(fit$var.coef)))))
```

```
##          ar1      ma1  intercept
## 0.000000000 0.004745202 0.000000000
```

Therefore, the The best model is ARMA(1,1). Now we will do forecasting and provide 95% prediction interval.

```
detach("package:itsmr")
library(forecast)
forecast(fit, 30)


##      Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
##  99          9.733373  8.836344 10.63040  8.361485 11.10526
## 100          9.560436  8.249647 10.87122  7.555758 11.56511
## 101          9.431615  7.939956 10.92327  7.150319 11.71291
## 102          9.335656  7.752526 10.91879  6.914467 11.75685
## 103          9.264177  7.632502 10.89585  6.768745 11.75961
## 104          9.210932  7.552934 10.86893  6.675242 11.74662
## 105          9.171270  7.498844 10.84370  6.613515 11.72902
## 106          9.141726  7.461348 10.82210  6.571810 11.71164
## 107          9.119718  7.434944 10.80449  6.543079 11.69636
## 108          9.103325  7.416116 10.79053  6.522962 11.68369
## 109          9.091113  7.402556 10.77967  6.508687 11.67354
```



```
## 110      9.082017  7.392711 10.77132  6.498447 11.66559
## 111      9.075241  7.385520 10.76496  6.491036 11.65945
## 112      9.070194  7.380243 10.76014  6.485637 11.65475
## 113      9.066434  7.376355 10.75651  6.481682 11.65119
## 114      9.063633  7.373484 10.75378  6.478773 11.64849
## 115      9.061547  7.371358 10.75174  6.476626 11.64647
## 116      9.059993  7.369783 10.75020  6.475039 11.64495
## 117      9.058836  7.368613 10.74906  6.473863 11.64381
## 118      9.057973  7.367744 10.74820  6.472990 11.64296
## 119      9.057331  7.367098 10.74756  6.472342 11.64232
## 120      9.056853  7.366617 10.74709  6.471861 11.64184
## 121      9.056496  7.366260 10.74673  6.471503 11.64149
## 122      9.056231  7.365994 10.74647  6.471236 11.64123
## 123      9.056033  7.365796 10.74627  6.471038 11.64103
## 124      9.055886  7.365648 10.74612  6.470890 11.64088
## 125      9.055776  7.365538 10.74601  6.470780 11.64077
## 126      9.055694  7.365456 10.74593  6.470699 11.64069
## 127      9.055633  7.365396 10.74587  6.470638 11.64063
## 128      9.055588  7.365350 10.74583  6.470592 11.64058
```

```
plot(forecast(fit, 30))
```

Forecasts from ARIMA(1,0,1) with non-zero mean



Out-of-sample forecasting error

Other than information criteria, out-of-sample forecasting error is widely used to determine the best model. In short, our best model is the one shows the smallest forecasting error. For example, 1-step out-of-sample forecasting error is given by

$$E(X_{t+1} - \hat{X}_{t+1})^2.$$

We will estimate this quantity by using cross-validation. In time series context, to preserve dependence structure, we divide entire sample data of size $N + m$ into two:

Training set X_1, \dots, X_N

Test set X_{N+1}, \dots, X_{N+m}

Then, we can approximate 1-step MSPE as

$$\frac{1}{m} \sum_{t=N}^{N+m-1} (X_{t+1} - \hat{X}_{t+1})^2.$$

Here are the procedure to calculate 1-step out-of-step forecasting error

1. Step1 : Find the best model using Training set X_1, \dots, X_N
2. Step2 : For $t = N, \dots, N + m - 1$ (corresponding to Test set): Find the 1-step ahead estimate X_{t+1} based on observation X_1, \dots, X_t based on the model parameters estimated in Step1. Or alternatively,
3. Step2' : We can re-estimate the model parameters in Step 2 using sample X_1, \dots, X_t .

In-class exercise

Implement and evaluate 1-step-out of sample forecasting error for mysterious.txt in HW5 for the last 30 observations. Which model would you select for this dataset?

```
## Out of sample forecasting practice
```

```
myst = scan("mysterious.txt");
```

```
library(forecast)
```

```
auto.arima(myst, d=0);
```

```
## Series: myst
```

```
## ARIMA(3,0,1) with zero mean
```

```
##
```

```
## Coefficients:
```

```
##          ar1      ar2      ar3      ma1
```

```
##      -0.1305  0.3945 -0.3174  0.8219
```

```
## s.e.   0.1078  0.0740  0.0683  0.0967
```

```
##
```

```
## sigma^2 estimated as 0.133: log likelihood=-80.37
```

```
## AIC=170.74  AICc=171.05  BIC=187.23
```

```
m=30; n = length(myst);
```

```
N = n-m;
```

```
testindex = (N+1):n;
```

```
p=3; q=1;
```

```
err = numeric(m);
```

```
for(i in 1:m){
```

```

trainindex = 1:(N+i-1);
fit = arima(myst[trainindex], order=c(p,0,q), include.mean=FALSE);
Xhat = forecast(fit, h=1)$mean;
err[i] = (myst[N+i] - Xhat)^2;
}
mean(err)

```

```
## [1] 0.1104433
```

```
## In-terms of out-of-sample forecasting ARMA(1,1) is better...
```

```

p=1; q=1;
err = numeric(m);
for(i in 1:m){
  trainindex = 1:(N+i-1);
  fit = arima(myst[trainindex], order=c(p,0,q), include.mean=FALSE);
  Xhat = forecast(fit, h=1)$mean;
  err[i] = (myst[N+i] - Xhat)^2;
}
mean(err)

```

```
## [1] 0.09931474
```