

```

> #####
> # Statistical Modelling & Machine Learning #
> # R Example3 #
> #####
>
> options(warn = -1) # Turn off warning message
>
> # Data: Breast Cancer Wisconsin Data
>
> dat = read.csv("wdbc.csv", header = F)
>
> x.name = c("radius", "texture", "perimeter", "area", "smoothness",
+           "compactness", "concavity", "concave_points", "symmetry",
+           "fractal_dimension")
>
> names(dat) = c("id", "diagnosis", paste0(x.name, "_mean"),
+               paste0(x.name, "_se"), paste0(x.name, "_worst"))
>
> dat = dat[, -1]
>
> head(dat)
diagnosis radius_mean texture_mean perimeter_mean area_mean smoothness_mean
1 M 17.99 10.38 122.80 1001.0 0.11840
2 M 20.57 17.77 132.90 1326.0 0.08474
3 M 19.69 21.25 130.00 1203.0 0.10960
4 M 11.42 20.38 77.58 386.1 0.14250
5 M 20.29 14.34 135.10 1297.0 0.10030
6 M 12.45 15.70 82.57 477.1 0.12780
compactness_mean concavity_mean concave_points_mean symmetry_mean
fractal_dimension_mean
1 0.27760 0.3001 0.14710 0.2419 0.07871
2 0.07864 0.0869 0.07017 0.1812 0.05667
3 0.15990 0.1974 0.12790 0.2069 0.05999
4 0.28390 0.2414 0.10520 0.2597 0.09744
5 0.13280 0.1980 0.10430 0.1809 0.05883
6 0.17000 0.1578 0.08089 0.2087 0.07613
radius_se texture_se perimeter_se area_se smoothness_se compactness_se concavity_se
1 1.0950 0.9053 8.589 153.40 0.006399 0.04904 0.05373
2 0.5435 0.7339 3.398 74.08 0.005225 0.01308 0.01860
3 0.7456 0.7869 4.585 94.03 0.006150 0.04006 0.03832
4 0.4956 1.1560 3.445 27.23 0.009110 0.07458 0.05661
5 0.7572 0.7813 5.438 94.44 0.011490 0.02461 0.05688
6 0.3345 0.8902 2.217 27.19 0.007510 0.03345 0.03672
concave_points_se symmetry_se fractal_dimension_se radius_worst texture_worst
1 0.01587 0.03003 0.006193 25.38 17.33
2 0.01340 0.01389 0.003532 24.99 23.41
3 0.02058 0.02250 0.004571 23.57 25.53
4 0.01867 0.05963 0.009208 14.91 26.50
5 0.01885 0.01756 0.005115 22.54 16.67
6 0.01137 0.02165 0.005082 15.47 23.75
perimeter_worst area_worst smoothness_worst compactness_worst concavity_worst
1 184.60 2019.0 0.1622 0.6656 0.7119
2 158.80 1956.0 0.1238 0.1866 0.2416
3 152.50 1709.0 0.1444 0.4245 0.4504
4 98.87 567.7 0.2098 0.8663 0.6869
5 152.20 1575.0 0.1374 0.2050 0.4000
6 103.40 741.6 0.1791 0.5249 0.5355
concave_points_worst symmetry_worst fractal_dimension_worst
1 0.2654 0.4601 0.11890
2 0.1860 0.2750 0.08902
3 0.2430 0.3613 0.08758
4 0.2575 0.6638 0.17300
5 0.1625 0.2364 0.07678
6 0.1741 0.3985 0.12440
>
> # Principal Component Analysis (PCA) -----
>
> pr = prcomp(dat[, 2:31], center = TRUE, scale = TRUE)
> summary(pr)

```

Importance of components:


	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172	0.69037	0.6457
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251	0.01589	0.0139
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010	0.92598	0.9399

	PC10	PC11	PC12	PC13	PC14	PC15	PC16	PC17
Standard deviation	0.59219	0.5421	0.51104	0.49128	0.39624	0.30681	0.28260	0.24372
Proportion of Variance	0.01169	0.0098	0.00871	0.00805	0.00523	0.00314	0.00266	0.00198
Cumulative Proportion	0.95157	0.9614	0.97007	0.97812	0.98335	0.98649	0.98915	0.99113

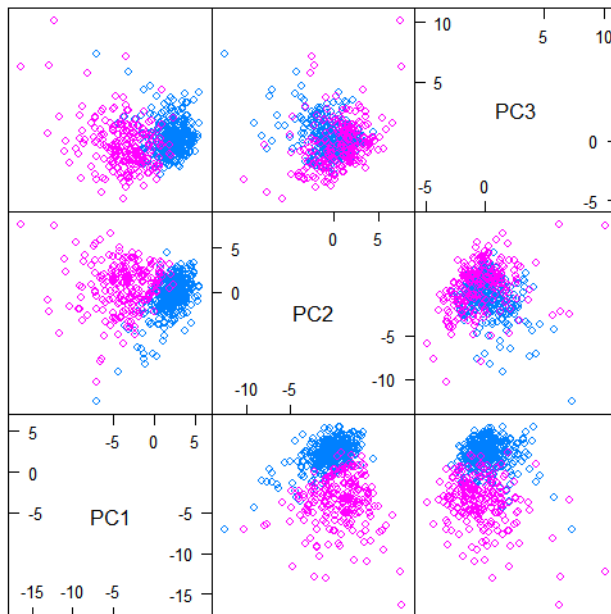
	PC18	PC19	PC20	PC21	PC22	PC23	PC24	PC25	PC26
Standard deviation	0.22939	0.22244	0.17652	0.1731	0.16565	0.15602	0.1344	0.12442	0.09043
Proportion of Variance	0.00175	0.00165	0.00104	0.0010	0.00091	0.00081	0.0006	0.00052	0.00027
Cumulative Proportion	0.99288	0.99453	0.99557	0.9966	0.99749	0.99830	0.9989	0.99942	0.99969

	PC27	PC28	PC29	PC30
Standard deviation	0.08307	0.03987	0.02736	0.01153
Proportion of Variance	0.00023	0.00005	0.00002	0.00000
Cumulative Proportion	0.99992	0.99997	1.00000	1.00000

```
>
> # Scree plot
> screeplot(pr, type = 'l', npcs = 15, main = 'Scree plot')
```

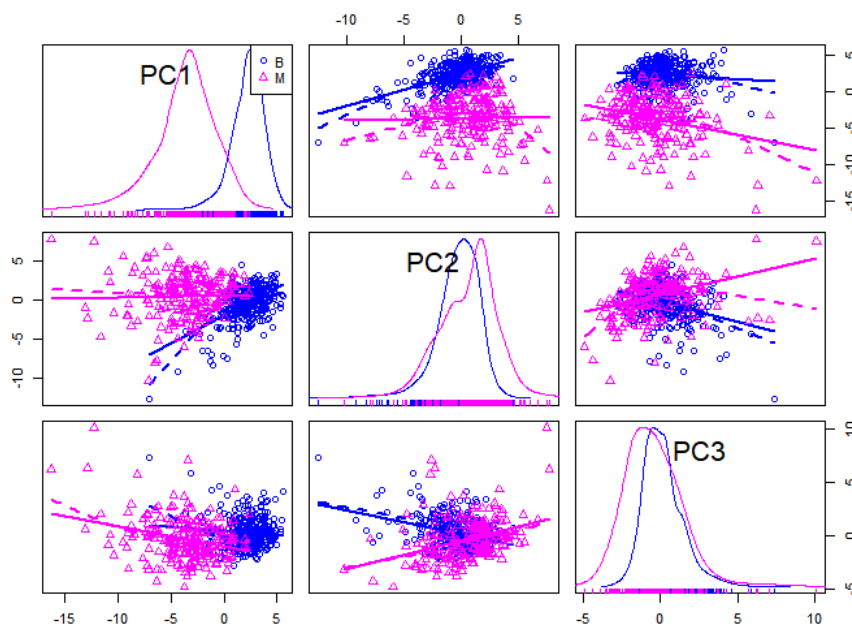


```
> # Visualization: Scatter plot matrix
>
> library(lattice)
>
> pc.dat = data.frame(type = dat$diagnosis, pr$x[, 1:3])
>
> pc.dat$type = as.factor(pc.dat$type)
>
> splom(~pc.dat[, 2:4], groups=type, data=pc.dat, panel=panel.superpose)
```



Scatter Plot Matrix

```
> install.packages('car')
> library(car)
> scatterplotMatrix(~PC1+PC2+PC3|type, data=pc.dat)
```




```
> # Application to Logistic regression:
> install.packages('boot')
> library(boot)
>
> dat$diagnosis = as.factor(dat$diagnosis)
>
> # Comparison of CV statistics:
> set.seed(1)
> fit = glm(diagnosis~., data=dat, family=binomial)
> cv.glm(dat, fit, K=5)$delta
[1] 0.05982303 0.82816919
>
```

```

> fit1 = glm(type~., data=pc.dat, family=binomial)
> cv.glm(pc.dat, fit1, K=5)$delta
[1] 0.03728052 0.03671780
>
> # Principal Curve -----
> install.packages('pncurve')
> library(pncurve)
>
> # Simple example
> set.seed(1)
> n=100
> x1 = runif(n, -1, 1)
> x2 = x1^2 + rnorm(n, sd = 0.1)
> z = cbind(x1, x2)
>
> cor(x1, x2)
[1] 0.08086294
>
> fit = principal_curve(z)
>
> plot(x1, x2)
> lines(sort(x1), (sort(x1))^2, col='red', lty=2)
> lines(fit, col='blue')
> whiskers(z, fit$s)



```





```

> # WDBC data application:
>
> fit = principal_curve(as.matrix(dat[, 2:31]))
>
> # Density of two groups in principal curve and PC1
> par(mfrow=c(1, 2))
> plot(density(fit$lambda[dat$diagnosis == 'B']), col='red',
+       xlim=c(500, 5500), main='Principal Curve')
> lines(density(fit$lambda[dat$diagnosis == 'M']), col='blue')
>
> plot(density(pc.dat[dat$diagnosis == 'B', 2]), col='red',
+       xlim=c(-15, 7), main='PC1')
> lines(density(pc.dat[dat$diagnosis == 'M', 2]), col='blue')

```



```
> # Density for principal curve and PC1
> par(mfrow=c(1, 2))
> plot(density(fit$lambda), col='red', main='Principal Curve')
> plot(density(pc.dat[, 2]), col='red', main='PC1')
```




```
>
> dat1 = cbind(dat, pcurve=fit$lambda)
> dat1$diagnosis = as.factor(dat1$diagnosis)
>
> fit2 = glm(diagnosis~pcurve, data=dat1, family=binomial)
> cv.glm(dat1, fit2, K=5)$delta
[1] 0.06816360 0.06812241
>
> # Projection onto Principal curve
> new.obs = as.matrix(dat[31:40, 2:31])
> project_to_curve(new.obs, fit2$lambda)
      31      32      33      34      35      36      37      38
4000.6148 2994.0477 1240.4169 1777.8542 2435.6018 2478.3771  0.0000 271.0701
      39      40
 496.7614 544.8027
> # arc-length along the curve.
```

```

> # Kernel PCA -----
> install.packages('kernelab')
> library(kernelab)
>
> x = dat[, 2:31]
> fit = kpca(~., data=x, kernel='rbfdot', kpar=list(sigma=3), features=2)
> # feature: # of PC's
>
> # Kernel PC's
> pc = pcv(fit)
>
> B = pc[dat$diagnosis=='B', ]
> M = pc[dat$diagnosis=='M', ]
>
> par(mfrow=c(1, 1))
> plot(B, col='red', xlab='KPC1', ylab='KPC2')
> points(M, col='blue')

```



```

>
> # New observations
> predict(fit, new.obs)
      [, 1]      [, 2]
31  1.672321e-08 -5.958310e-08
32 -2.204971e-07  4.754435e-08
33 -3.603358e-08 -4.478933e-08
34 -3.231095e-07 -6.363222e-08
35  5.273944e-08 -5.169283e-08
36  3.538922e-08 -5.199263e-08
37  4.219275e-08 -7.082468e-08
38  1.504703e-07 -4.966250e-08
39 -9.474597e-07 -5.998169e-08
40 -3.543559e-06 -8.392185e-08

```

```

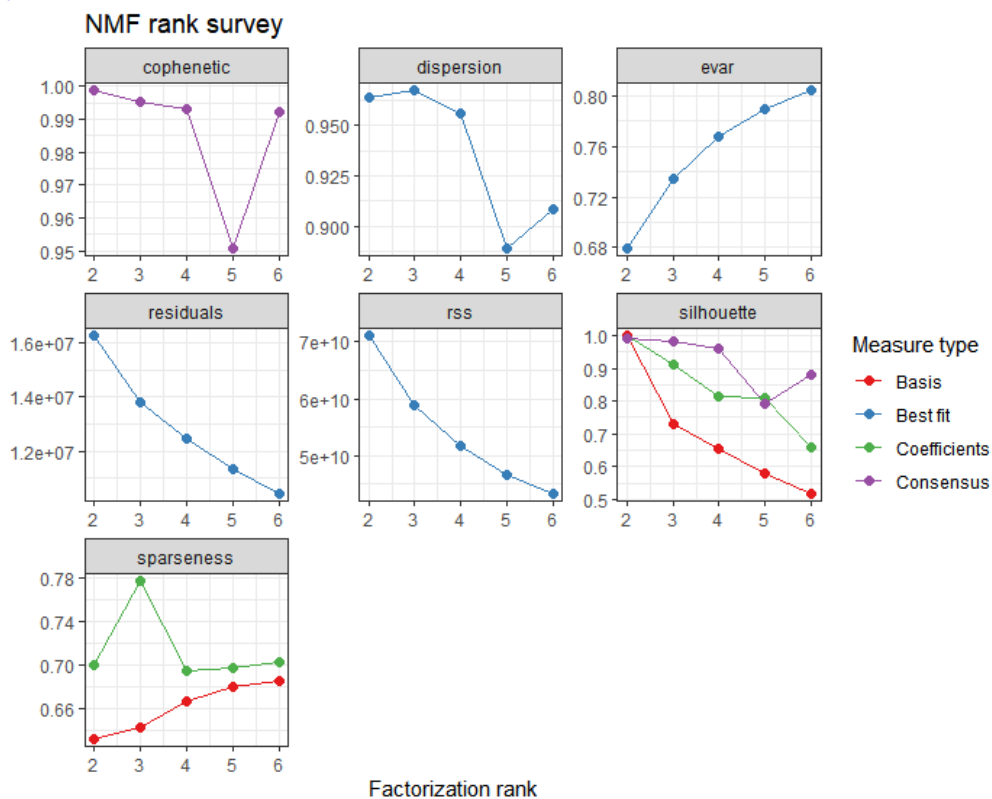
> # Non-negative matrix factorization -----
>
> if (!requireNamespace("BiocManager", quietly = TRUE))
+   install.packages("BiocManager")
> BiocManager::install("Biobase")
>
> install.packages('NMF')
> library(NMF)
>
> ?esGolub
>
> data(esGolub)
>

```

```

> dim(esGolub)
Features Samples
5000         38
>
> res = nmf(esGolub, rank = 3, seed=123456)
>
> # W matrix
> W = basis(res)
> dim(W)
[1] 5000    3
>
> # H matrix
> H = coef(res)
> dim(H)
[1] 3 38
>
>
> if(requireNamespace("Biobase", quietly=TRUE))
+ {
+   estim.r = nmf(esGolub, 2:6, nrun=10, seed=123456)
+   plot(estim.r)
+ }

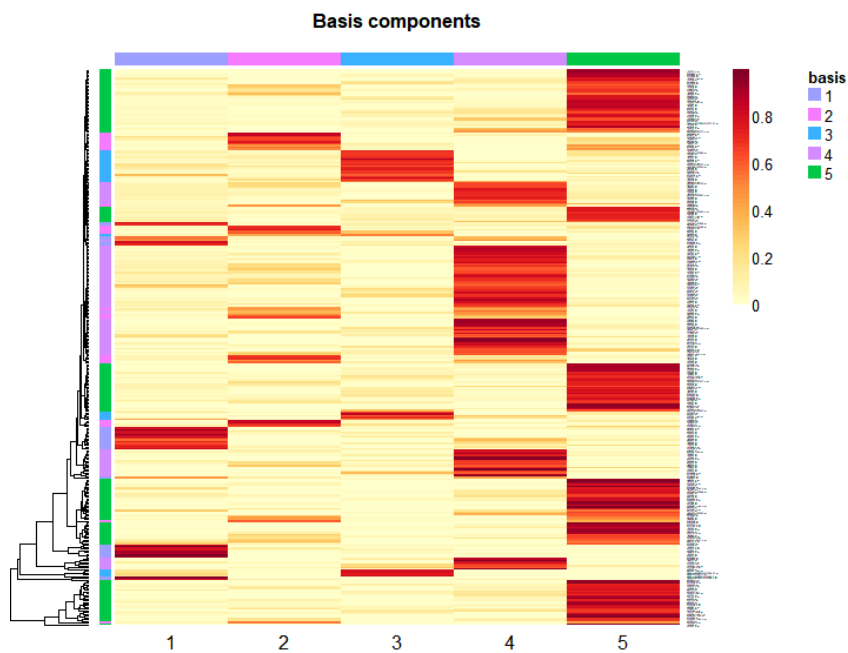
```



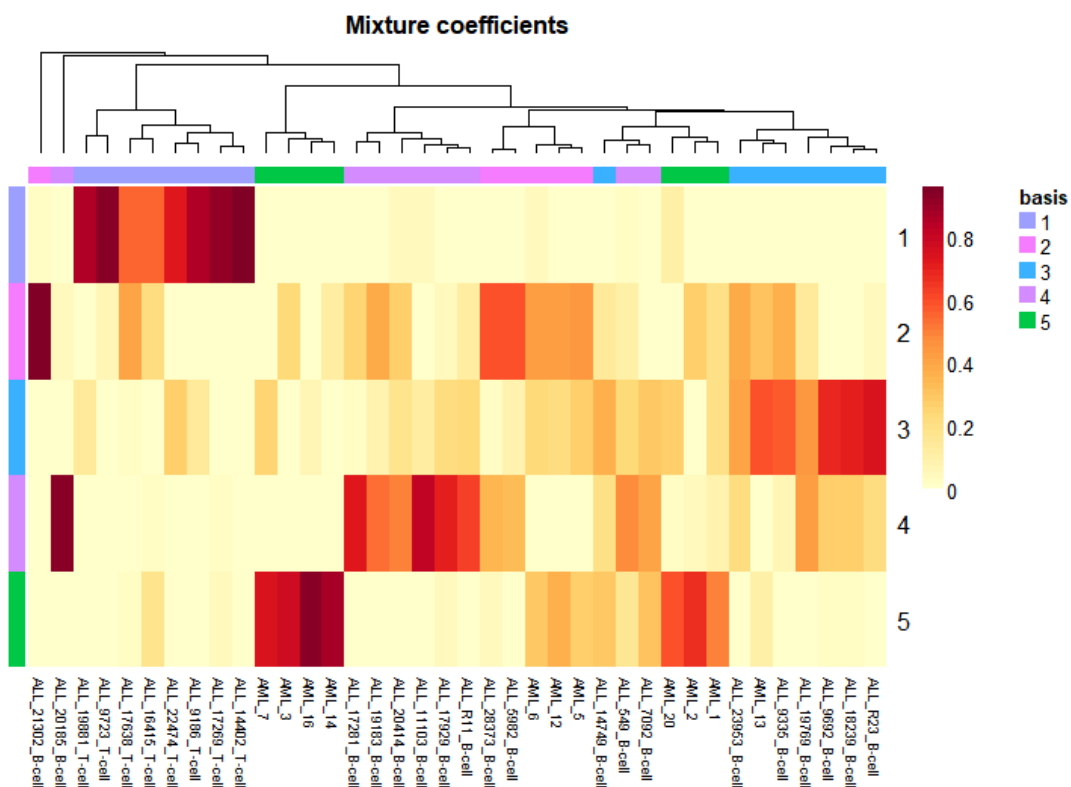
```

>
> res = nmf(esGolub, rank = 5, seed=123456)
>
> # Visualization
> basismap(res, subsetRow=TRUE)

```



```
>
> coefmap(res)
```



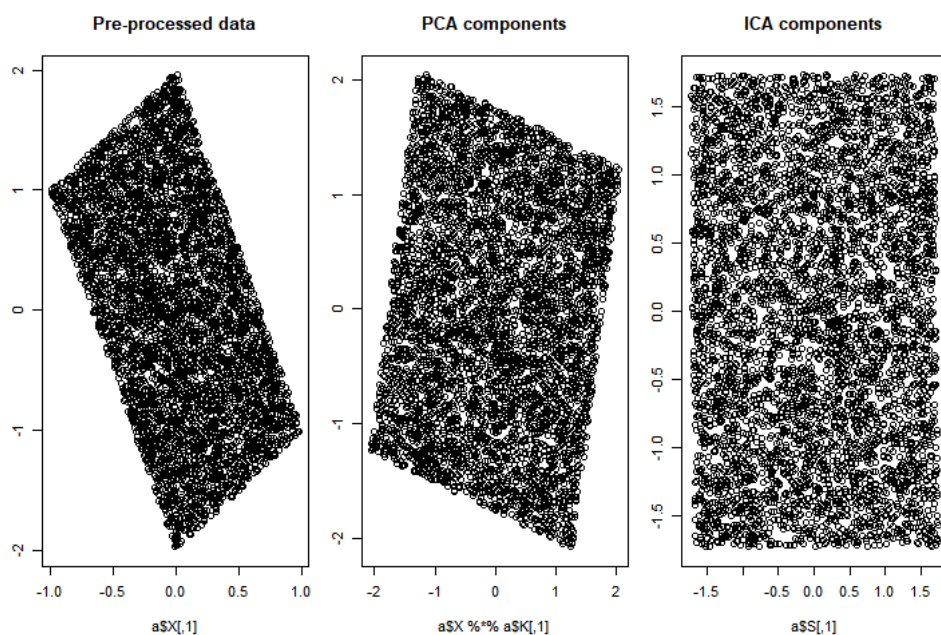
```
>
> # Independent component analysis -----
>
> install.packages('fastICA')
> library(fastICA)
>
>
> # Ex1:
>
> S = matrix(runif(10000), 5000, 2)
```



```

> A = matrix(c(1, 1, -1, 3), 2, 2, byrow = TRUE)
> X = S %*% A
> a = fastICA(X, 2, alg.typ = "parallel", fun = "logcosh", alpha = 1,
+           method = "C", row.norm = FALSE, maxi t = 200,
+           tol = 0.0001, verbose = TRUE)
Centering
Whitening
Symmetric FastICA using logcosh approx. to neg-entropy function
Iteration 1 tol=0.104526
Iteration 2 tol=0.003510
Iteration 3 tol=0.000001
> par(mfrow = c(1, 3))
> plot(a$X, main = "Pre-processed data")
> plot(a$X %*% a$K, main = "PCA components")
> plot(a$S, main = "ICA components")

```



```

> # Ex2:
>
> S = cbind(sin((1:1000)/20), rep((((1:200)-100)/100), 5))
> A = matrix(c(0.291, 0.6557, -0.5439, 0.5572), 2, 2)
> X = S %*% A
> a = fastICA(X, 2, alg.typ = "parallel", fun = "logcosh", alpha = 1,
+           method = "R", row.norm = FALSE, maxi t = 200,
+           tol = 0.0001, verbose = TRUE)
Centering
Whitening
Symmetric FastICA using logcosh approx. to neg-entropy function
Iteration 1 tol = 0.004695925
Iteration 2 tol = 4.449089e-06
>
> par(mfcol = c(2, 3))
> plot(1:1000, S[, 1], type = "l", main = "Original Signals",
+      xlab = "", ylab = "")
> plot(1:1000, S[, 2], type = "l", xlab = "", ylab = "")
> plot(1:1000, X[, 1], type = "l", main = "Mixed Signals",
+      xlab = "", ylab = "")
> plot(1:1000, X[, 2], type = "l", xlab = "", ylab = "")
> plot(1:1000, a$S[, 1], type = "l", main = "ICA source estimates",
+      xlab = "", ylab = "")
> plot(1:1000, a$S[, 2], type = "l", xlab = "", ylab = "")

```

